

```
In [16]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import OneHotEncoder, PolynomialFeatures
from sklearn.metrics import r2_score
```

```
In [3]: data=pd.read_csv("C:\\Users\\Lenovo\\Desktop\\CodeSoft\\Task4_SalesPrediction\\adve
data.head(10)
```

```
Out[3]:
```

|   | TV    | Radio | Newspaper | Sales |
|---|-------|-------|-----------|-------|
| 0 | 230.1 | 37.8  | 69.2      | 22.1  |
| 1 | 44.5  | 39.3  | 45.1      | 10.4  |
| 2 | 17.2  | 45.9  | 69.3      | 12.0  |
| 3 | 151.5 | 41.3  | 58.5      | 16.5  |
| 4 | 180.8 | 10.8  | 58.4      | 17.9  |
| 5 | 8.7   | 48.9  | 75.0      | 7.2   |
| 6 | 57.5  | 32.8  | 23.5      | 11.8  |
| 7 | 120.2 | 19.6  | 11.6      | 13.2  |
| 8 | 8.6   | 2.1   | 1.0       | 4.8   |
| 9 | 199.8 | 2.6   | 21.2      | 15.6  |

```
In [6]: data.shape
```

```
Out[6]: (200, 4)
```

```
In [7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV          200 non-null    float64
1    Radio        200 non-null    float64
2    Newspaper    200 non-null    float64
3    Sales        200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
In [8]: data.describe()
```


```
Out[8]:
```

|              | TV         | Radio      | Newspaper  | Sales      |
|--------------|------------|------------|------------|------------|
| <b>count</b> | 200.000000 | 200.000000 | 200.000000 | 200.000000 |
| <b>mean</b>  | 147.042500 | 23.264000  | 30.554000  | 15.130500  |
| <b>std</b>   | 85.854236  | 14.846809  | 21.778621  | 5.283892   |
| <b>min</b>   | 0.700000   | 0.000000   | 0.300000   | 1.600000   |
| <b>25%</b>   | 74.375000  | 9.975000   | 12.750000  | 11.000000  |
| <b>50%</b>   | 149.750000 | 22.900000  | 25.750000  | 16.000000  |
| <b>75%</b>   | 218.825000 | 36.525000  | 45.100000  | 19.050000  |
| <b>max</b>   | 296.400000 | 49.600000  | 114.000000 | 27.000000  |


```
In [10]: data.isnull().sum()*100/data.shape[0]
```

```
Out[10]: TV          0.0
Radio       0.0
Newspaper   0.0
Sales       0.0
dtype: float64
```

```
In [17]: fig, axs = plt.subplots(3, figsize = (5, 5))
plt1 = snr.boxplot(data['TV'], ax = axs[0])
plt2 = snr.boxplot(data['Newspaper'], ax = axs[1])
plt3 = snr.boxplot(data['Radio'], ax = axs[2])
plt.tight_layout()
```




```
In [18]: snr.boxplot(data['Sales'])
plt.show()
```




```
In [21]: snr.pairplot(data, x_vars=['TV', 'Newspaper', 'Radio'], y_vars='Sales', height=4, a
plt.show()
```

C:\Users\Lenovo\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning:  
The figure layout has changed to tight  
self.\_figure.tight\_layout(\*args, \*\*kwargs)



```
In [22]: sns.heatmap(data.corr(), cmap="YlGnBu", annot = True)
plt.show()
```



```
In [23]: X = data['TV']
        Y = data['Sales']
```

```
In [25]: from sklearn.model_selection import train_test_split
        X_train, X_test, Y_train, Y_test = train_test_split(X, Y, train_size = 0.7, test_si
```

```
In [26]: X_train.head()
```

```
Out[26]: 74    213.4
         3    151.5
        185   205.0
         26   142.9
         90   134.3
        Name: TV, dtype: float64
```

```
In [27]: Y_train.head()
```

```
Out[27]: 74    17.0
         3    16.5
        185    22.6
         26    15.0
         90    14.0
        Name: Sales, dtype: float64
```

```
In [28]: import statsmodels.api as sm
```

```
In [30]: X_train_sm = sm.add_constant(X_train)
        lr = sm.OLS(Y_train, X_train_sm).fit()
```

```
In [31]: lr.params
```

```
Out[31]: const    6.948683
        TV      0.054546
        dtype: float64
```

```
In [32]: print(lr.summary())
```


# OLS Regression Results

|                   |                  |                     |          |
|-------------------|------------------|---------------------|----------|
| Dep. Variable:    | Sales            | R-squared:          | 0.816    |
| Model:            | OLS              | Adj. R-squared:     | 0.814    |
| Method:           | Least Squares    | F-statistic:        | 611.2    |
| Date:             | Fri, 20 Dec 2024 | Prob (F-statistic): | 1.52e-52 |
| Time:             | 13:08:46         | Log-Likelihood:     | -321.12  |
| No. Observations: | 140              | AIC:                | 646.2    |
| Df Residuals:     | 138              | BIC:                | 652.1    |
| Df Model:         | 1                |                     |          |
| Covariance Type:  | nonrobust        |                     |          |
| =====             |                  |                     |          |
|                   | coef             | std err             | t        |
|                   |                  |                     | P> t     |
|                   |                  |                     | [0.025   |
|                   |                  |                     | 0.975]   |
| -----             |                  |                     |          |
| const             | 6.9487           | 0.385               | 18.068   |
| TV                | 0.0545           | 0.002               | 24.722   |
|                   |                  |                     | 0.000    |
|                   |                  |                     | 0.000    |
|                   |                  |                     | 0.050    |
|                   |                  |                     | 0.059    |
| =====             |                  |                     |          |
| Omnibus:          | 0.027            | Durbin-Watson:      | 2.196    |
| Prob(Omnibus):    | 0.987            | Jarque-Bera (JB):   | 0.150    |
| Skew:             | -0.006           | Prob(JB):           | 0.928    |
| Kurtosis:         | 2.840            | Cond. No.           | 328.     |
| =====             |                  |                     |          |

## Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [33]: plt.scatter(X_train, Y_train)
plt.plot(X_train, 6.948 + 0.054*X_train, 'r')
plt.show()
```



```
In [34]: Y_train_pred = lr.predict(X_train_sm)
res = (Y_train - Y_train_pred)
```

```
In [35]: fig = plt.figure()
snr.distplot(res, bins = 15)
fig.suptitle('Error Terms', fontsize = 15)
```

```
plt.xlabel('Y_train - Y_train_pred', fontsize = 15)
plt.show()
```

C:\Users\Lenovo\AppData\Local\Temp\ipykernel\_6508\3636744094.py:2: UserWarning:


`distplot` is a deprecated function and will be removed in seaborn v0.14.0.

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
snr.distplot(res, bins = 15)
```

## Error Terms



```
In [36]: plt.scatter(X_train, res)
plt.show()
```



```
In [37]: X_test_sm = sm.add_constant(X_test)
Y_pred = lr.predict(X_test_sm)
```

```
In [38]: Y_pred.head()
```

```
Out[38]: 126    7.374140
104    19.941482
99     14.323269
92     18.823294
111    20.132392
dtype: float64
```

```
In [40]: from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
```


```
In [41]: np.sqrt(mean_squared_error(Y_test, Y_pred))
```

```
Out[41]: 2.019296008966232
```

```
In [42]: r_squared = r2_score(Y_test, Y_pred)
r_squared
```

```
Out[42]: 0.7921031601245659
```

```
In [43]: plt.scatter(X_test, Y_test)
plt.plot(X_test, 6.948 + 0.054 * X_test, 'r')
plt.show()
```



In [ ]: